

# DATA2002

## Decision trees and random forests

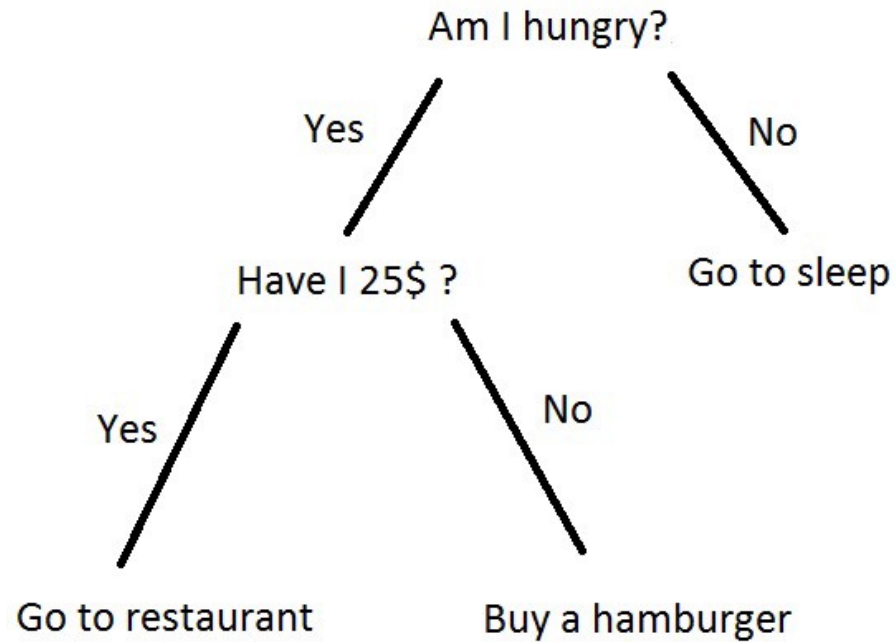
Garth Tarr



Decision trees

Random forests

# Decision trees



A decision tree determines the predicted outcome based on series of questions and conditions.

When making Decision Trees, there are several factors we take into consideration.

- What features do we make our decisions on?
- What is the threshold for classifying each question into a **yes** or **no** answer?

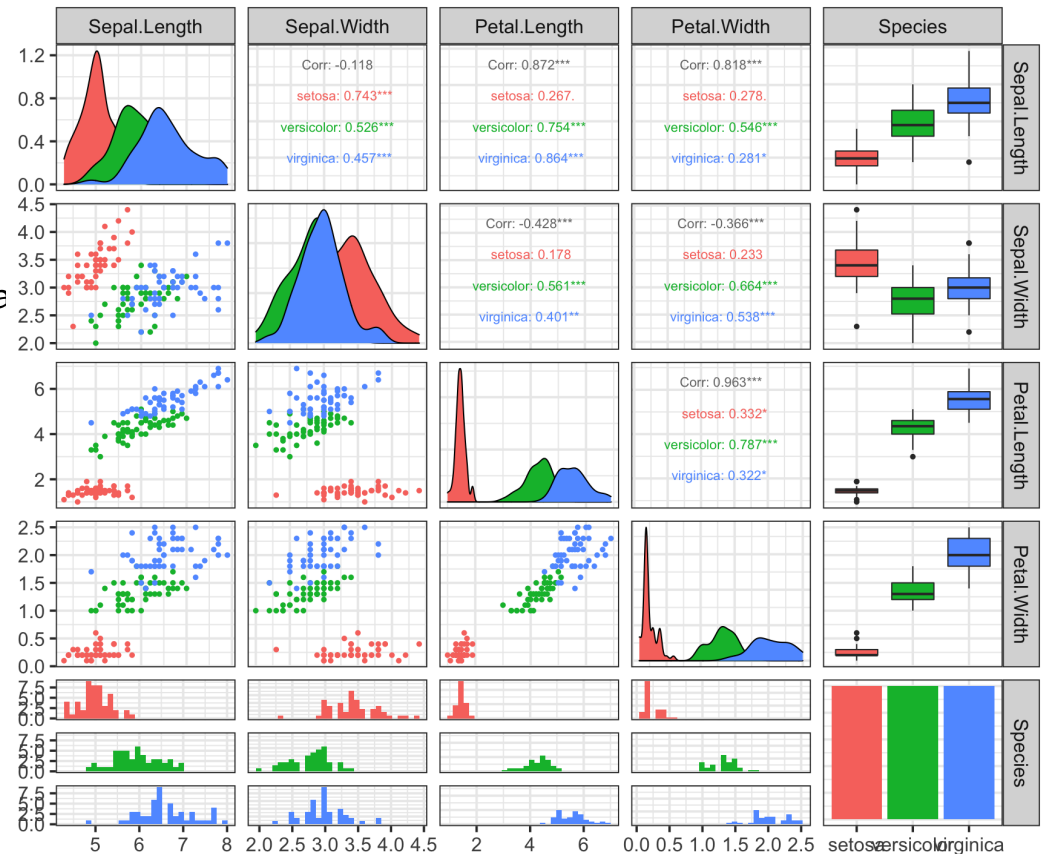
# Consider the iris dataset

```
library(tidyverse)
glimpse(iris)
```

```
## Rows: 150
## Columns: 5
## $ Sepal.Length <dbl> 5.1, 4.9, 4.7, 4.6, 5.0, 5.4,
## $ Sepal.Width <dbl> 3.5, 3.0, 3.2, 3.1, 3.6, 3.9,
## $ Petal.Length <dbl> 1.4, 1.4, 1.3, 1.5, 1.4, 1.7,
## $ Petal.Width <dbl> 0.2, 0.2, 0.2, 0.2, 0.2, 0.4,
## $ Species <fct> setosa, setosa, setosa, setosa
```

How can we construct a **decision tree** to determine the Species of iris given petal and sepal measurements?

```
library(GGally)
ggpairs(iris, mapping = aes(col = Species)) + t
```





# Trees in R

- We can use the **rpart** package (recursive partitioning) to build our decision tree.
- It uses a formula structure, much like `lm()` and `glm()` to identify the dependent variable (what we're trying to predict) and the explanatory variables (what information do we have available to help prediction).

```
library(rpart)
```

```
tree = rpart(Species ~ ., data = iris, method = "class")
```

- The formula `Species ~ .` in the code above says we want to predict `Species` using all the variables in the data frame `iris`.
- Specifying `method = "class"` means we want to build a classification (decision) tree.

```
tree
```

```
## n= 150
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
## 1) root 150 100 setosa (0.33333333 0.33333333 0.33333333)
##   2) Petal.Length< 2.45 50   0 setosa (1.00000000 0.00000000 0.00000000) *
##   3) Petal.Length>=2.45 100  50 versicolor (0.00000000 0.50000000 0.50000000)
##     6) Petal.Width< 1.75 54   5 versicolor (0.00000000 0.90740741 0.09259259) *
##     7) Petal.Width>=1.75 46   1 virginica (0.00000000 0.02173913 0.97826087) *
```

```
summary(tree)
```


```
## Call:
## rpart(formula = Species ~ ., data = iris, method = "class")
##      n= 150
##
##      CP nsplit rel error xerror      xstd
## 1 0.50      0      1.00   1.16 0.05127703
## 2 0.44      1      0.50   0.72 0.06118822
```

# Trees in R

```
library(rpart)
tree = rpart(
  Species ~ .,
  data = iris,
  method = "class")
```

```
library(rpart.plot)
rpart.plot(tree)
```

- Using the `rpart.plot()` function from the **rpart.plot** library, we can visualise the results of our classification tree.
- Each node shows the predicted class, the predicted probability of each class, and the percentage of observations in each node.






# Is it magic?

How does it work? In this tree, we only need to consider two variables.

```
p1 = iris %>% ggplot() +  
  aes(x = Petal.Length,  
      y = Petal.Width,  
      colour = Species) +  
  geom_point(size = 4) +  
  theme_classic(base_size = 30)
```

p1




# Is it magic?

The first branch is done to "best" split the data to create the most "pure" (homogenous) partitions.

`Petal.Length < 2.45`

```
p2 = p1 +  
  geom_vline(  
    aes(xintercept=2.45),  
    size = 2  
  )
```

p2



# Is it magic?


The next branch applies to observations that have

`Petal.Length > 2.45` and it tries to find the next best split of the data.

`Petal.width < 1.75`

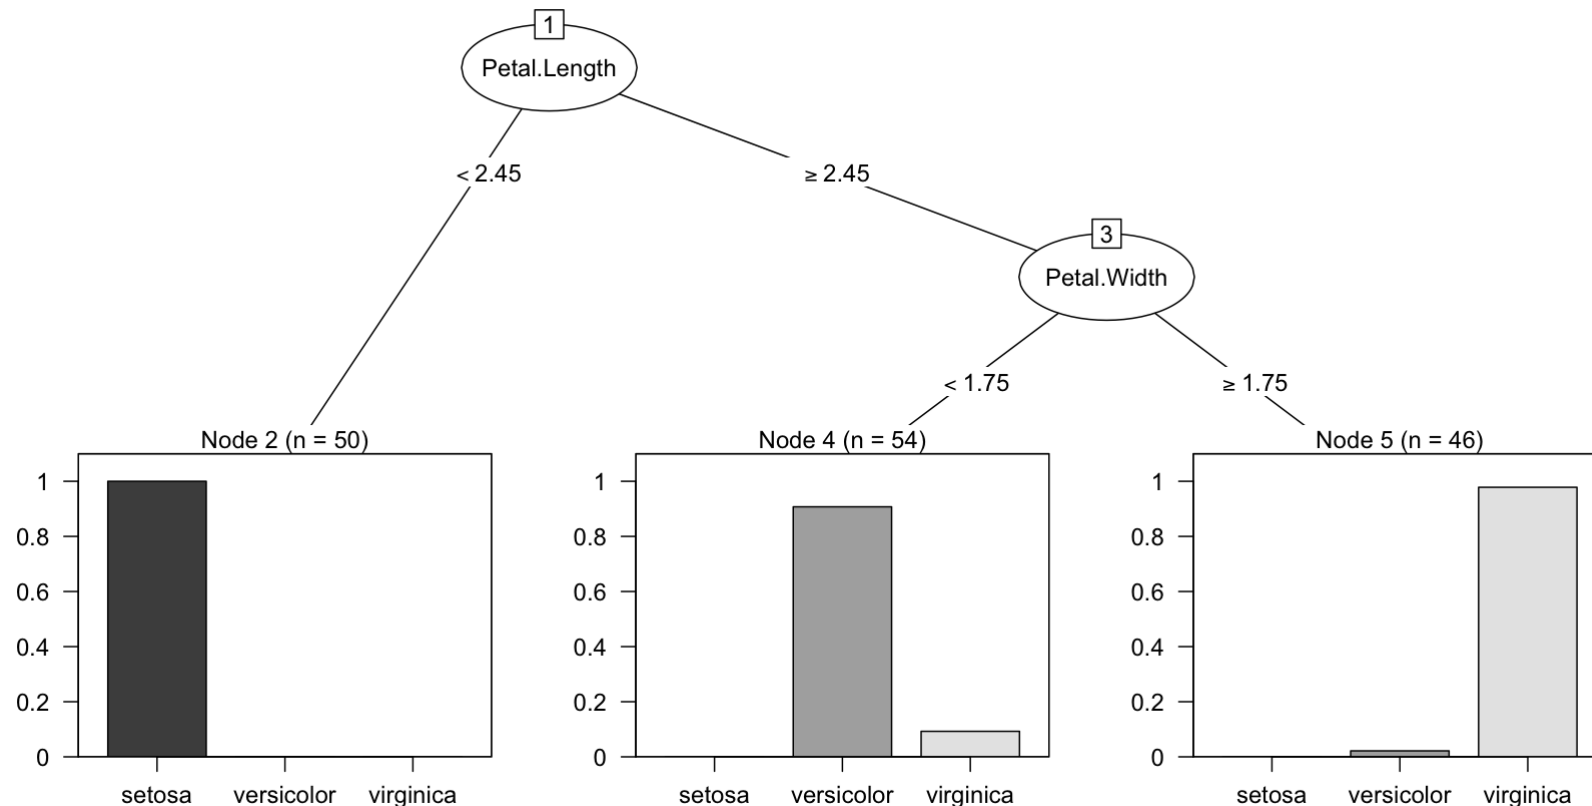
```
p3 = p2 +  
  geom_segment(  
    aes(x = 2.45, y = 1.75,  
        xend = 6.9, yend = 1.75),  
    size = 2,  
    colour = "black"  
  )
```

p3



# Alternative visualisation with partykit

```
# install.packages("partykit") # A Toolkit for Recursive Partytioning  
library(partykit)  
plot(as.party(tree))
```




# Making a prediction

How would we predict the species of a flower with

- Sepal.Length = 5.0
- Sepal.Width = 3.9
- Petal.Length = 1.4
- Petal.Width = 0.3

How about:

- Sepal.Length = 5.0
- Sepal.Width = 3.9
- Petal.Length = 3.4
- Petal.Width = 0.3



# Making a prediction

This was our fitted model:

```
tree <- rpart(Species ~ ., data = iris, method = "class")
```

Using `predict()`, we can predict the class of a new data point, much like for `lm()` and `glm()` objects.

```
new_data = data.frame(Sepal.Length = c(5.0, 5.0),
                      Sepal.Width = c(3.9, 3.9),
                      Petal.Length = c(1.4, 3.4),
                      Petal.Width = c(0.3, 0.3))

new_data
```

```
##   Sepal.Length Sepal.Width Petal.Length Petal.Width
## 1           5         3.9         1.4         0.3
## 2           5         3.9         3.4         0.3
```

```
predict(tree, new_data, type = "class")
```

```
##           1           2
##   setosa versicolor
## Levels: setosa versicolor virginica
```

# Assessing in sample performance in the iris data

```
library(caret)
predicted_species = predict(tree, type = "class")
confusionMatrix(
  data = predicted_species,
  reference = iris$Species)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##              Reference
## Prediction  setosa versicolor virginica
##   setosa      50         0         0
## versicolor    0         49         5
## virginica     0         1        45
```

```
##
```

```
## Overall Statistics
```

```
##
```

```
##              Accuracy : 0.96
##              95% CI : (0.915, 0.9852)
## No Information Rate : 0.3333
## P-Value [Acc > NIR] : < 2.2e-16
```

```
##
```

```
##              Kappa : 0.94
```

```
##
```

# Model selection

- How did our tree know to know to only use two splits?
- There is a **complexity parameter** that can be used to determine if a proposed new split *sufficiently* improves the predictive power or not.
- A choice is made whether to keep or "prune" a proposed new branch.
- Default is that a branch should decrease the error by 1%.
- This helps to avoid **overfitting**.

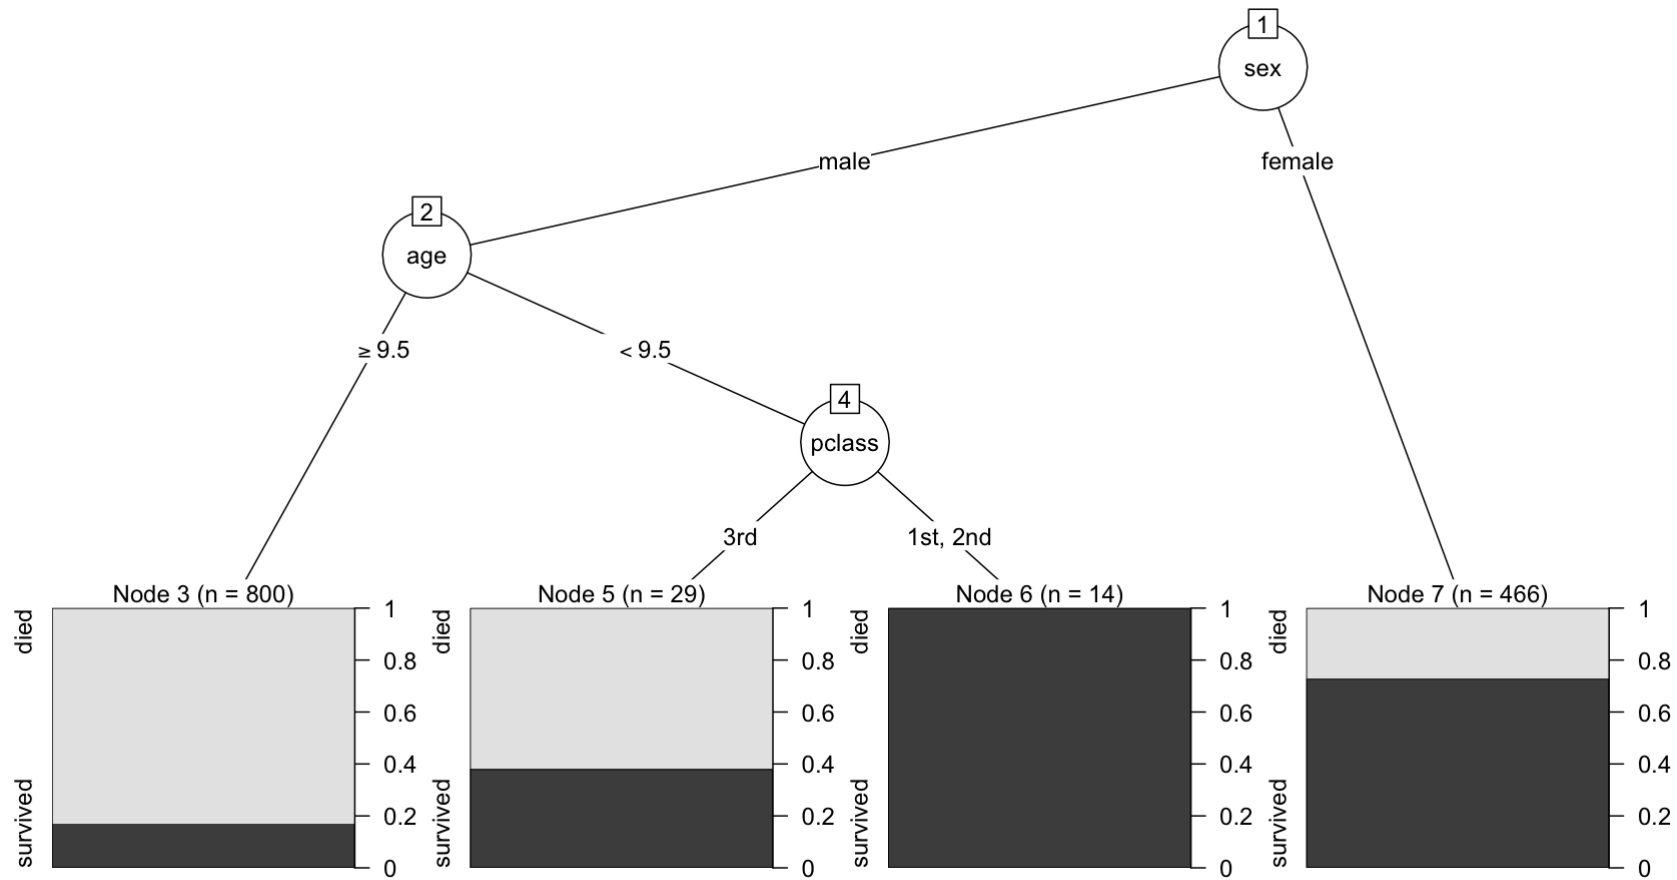


# Titanic tree

```
data("Titanicp", package = "vcdExtra")
titanic_tree = rpart(survived ~ sex + age + pclass, data = Titanicp, method = "class")
titanic_tree
```

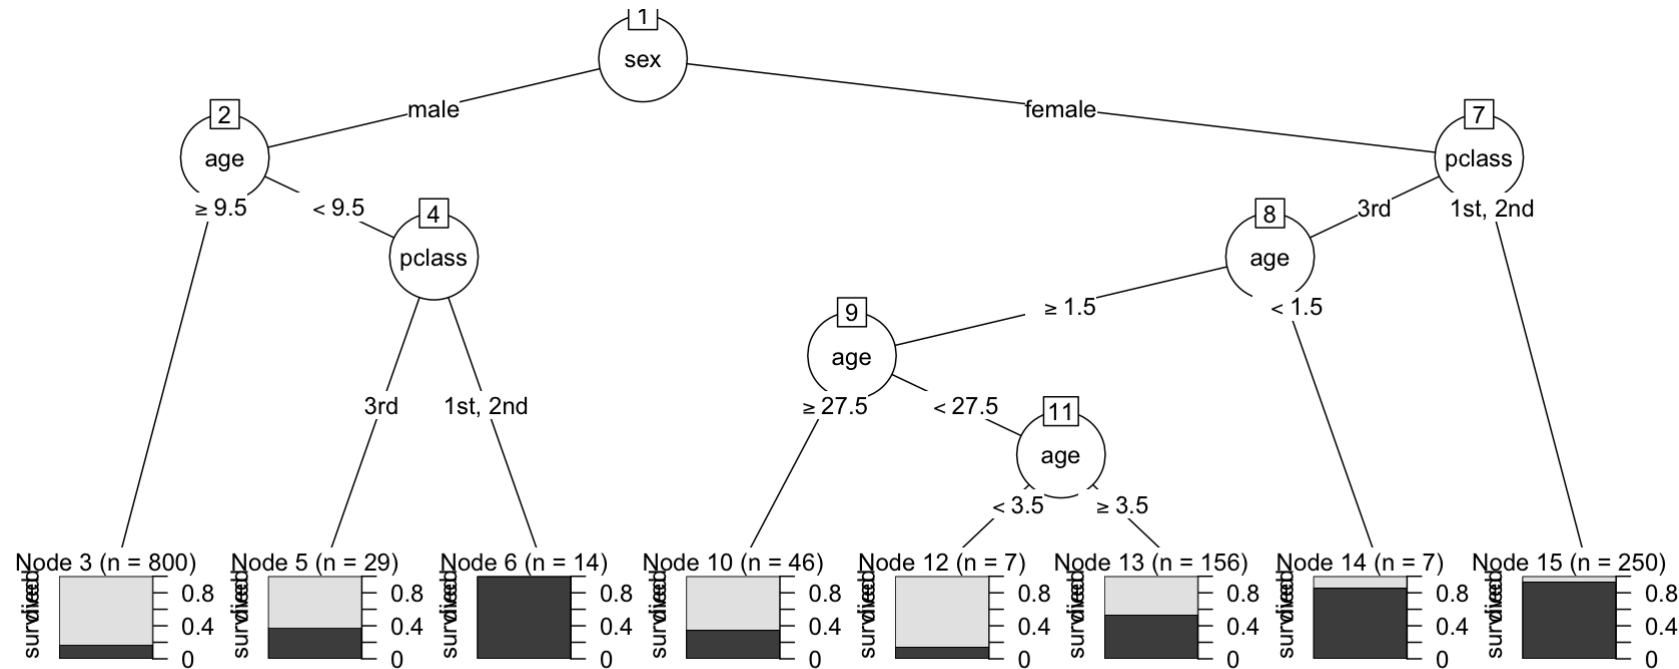
```
## n= 1309
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
## 1) root 1309 500 died (0.6180290 0.3819710)
##    2) sex=male 843 161 died (0.8090154 0.1909846)
##      4) age>=9.5 800 136 died (0.8300000 0.1700000) *
##      5) age< 9.5 43  18 survived (0.4186047 0.5813953)
##        10) pclass=3rd 29  11 died (0.6206897 0.3793103) *
##        11) pclass=1st,2nd 14  0 survived (0.0000000 1.0000000) *
##    3) sex=female 466 127 survived (0.2725322 0.7274678) *
```

```
plot(as.party(titanic_tree))
```




What if we lower the complexity parameter threshold, so that each new branch only needs to decrease the error by 0.9%?

```
titanic_tree0.9 = rpart(survived ~ sex + age + pclass, data = Titanicp, method = "class",  
                        control = rpart.control(cp = 0.009))  
plot(as.party(titanic_tree0.9))
```



What if we increase the complexity parameter threshold, so that each new branch needs to decrease the error by 2%?

```
titanic_tree2 = rpart(survived ~ sex + age + pclass, data = Titanicp, method = "class",  
                      control = rpart.control(cp = 0.02))  
plot(as.party(titanic_tree2))
```



# Evaluating (in-sample) performance

1% (default)

```
titanic_1_pred = predict(titanic_tree, type = "class")  
confusionMatrix(data=titanic_1_pred, reference = Titanicp$survived)$table
```

```
##           Reference  
## Prediction died survived  
##   died      682      147  
##   survived  127      353
```

```
confusionMatrix(data=titanic_1_pred, reference = Titanicp$survived)$overall[1]
```

```
## Accuracy  
## 0.7906799
```

# Evaluating (in-sample) performance

0.9%

```
titanic_0.9_pred = predict(titanic_tree0.9, type = "class")
confusionMatrix(data=titanic_0.9_pred,
                  reference = Titanicp$survived)$table
```

```
##           Reference
## Prediction died survived
##   died       718     164
##   survived   91     336
```

```
confusionMatrix(data=titanic_0.9_pred,
                  reference = Titanicp$survived)$overall[1]
```

```
## Accuracy
## 0.8051948
```

# Evaluating (in-sample) performance

2%

```
titanic_2_pred = predict(titanic_tree2, type = "class")
confusionMatrix(data=titanic_2_pred,
                  reference = Titanicp$survived)$table
```

```
##           Reference
## Prediction died survived
##   died      682      161
##   survived  127      339
```

```
confusionMatrix(data=titanic_2_pred,
                  reference = Titanicp$survived)$overall[1]
```

```
## Accuracy
## 0.7799847
```

# Performance benchmarking

```
table(Titanicp$survived)
```

```
##  
##      died survived  
##      809       500
```

What if our prediction model was just that everyone died? The accuracy would be:

```
809 / (809+500)
```

```
## [1] 0.618029
```

When considering performance, we should take into account that a "null" model might appear to give quite good performance when we have unbalanced group sizes.



# Evaluating (out-of-sample) performance

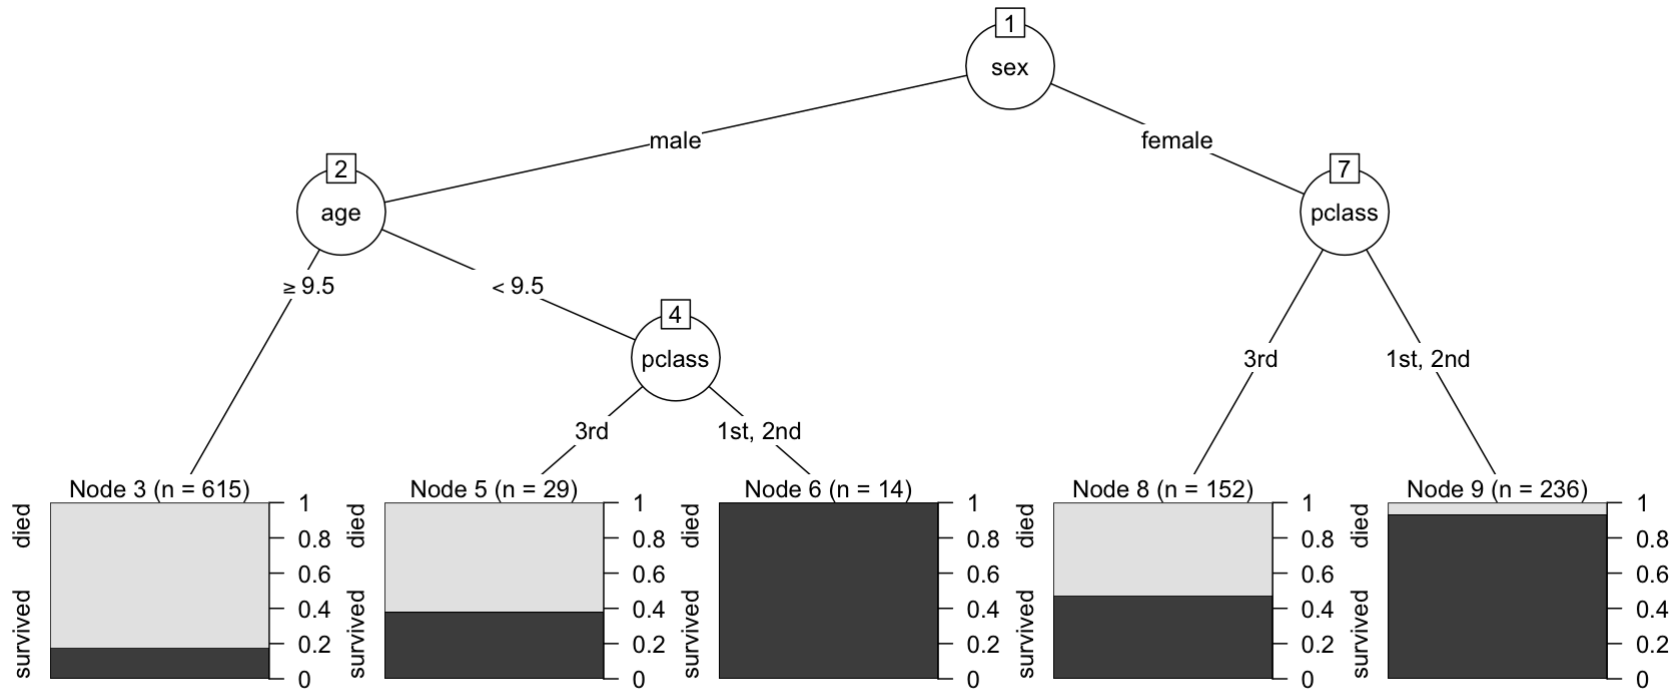
```
titanic_complete = Titanicp %>% select(survived, sex, age, pclass) %>% drop_na()
train(survived ~ sex + age + pclass, data = titanic_complete,
      method = "rpart", trControl = trainControl(method = "cv", number = 10))
```

```
## CART
##
## 1046 samples
##    3 predictor
##    2 classes: 'died', 'survived'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 942, 941, 942, 941, 941, 941, ...
## Resampling results across tuning parameters:
##
##    cp          Accuracy    Kappa
##    0.01639344  0.7696021    0.5002840
##    0.01873536  0.7648311    0.4910097
##    0.45901639  0.6748002    0.2449885
##
## Accuracy was used to select the optimal model using
## the largest value.
## The final value used for the model was cp = 0.01639344.
```

# Final model



The CV procedure suggested 1.6% for the complexity parameter.

```
titanic_final = rpart(survived ~ sex + age + pclass, data = titanic_complete,  
                      control = rpart.control(cp = 0.016))  
plot(as.party(titanic_final))
```




# Decision tree weaknesses

- Decision trees can become very complex very quickly - without a complexity penalty, it will happily continue until perfect classification (likely massively overfitting the data).
- The selected tree might be very sensitive to the complexity penalty
- Can only make decisions parallel to axes:



# A single tree is prone to overfitting



When you **overfit** you're modelling noise rather than signal.

# Random forests


# Random forests

In a **random forest** we grow many trees. Each one learns from different (sub)samples of observations and different combinations of variables.

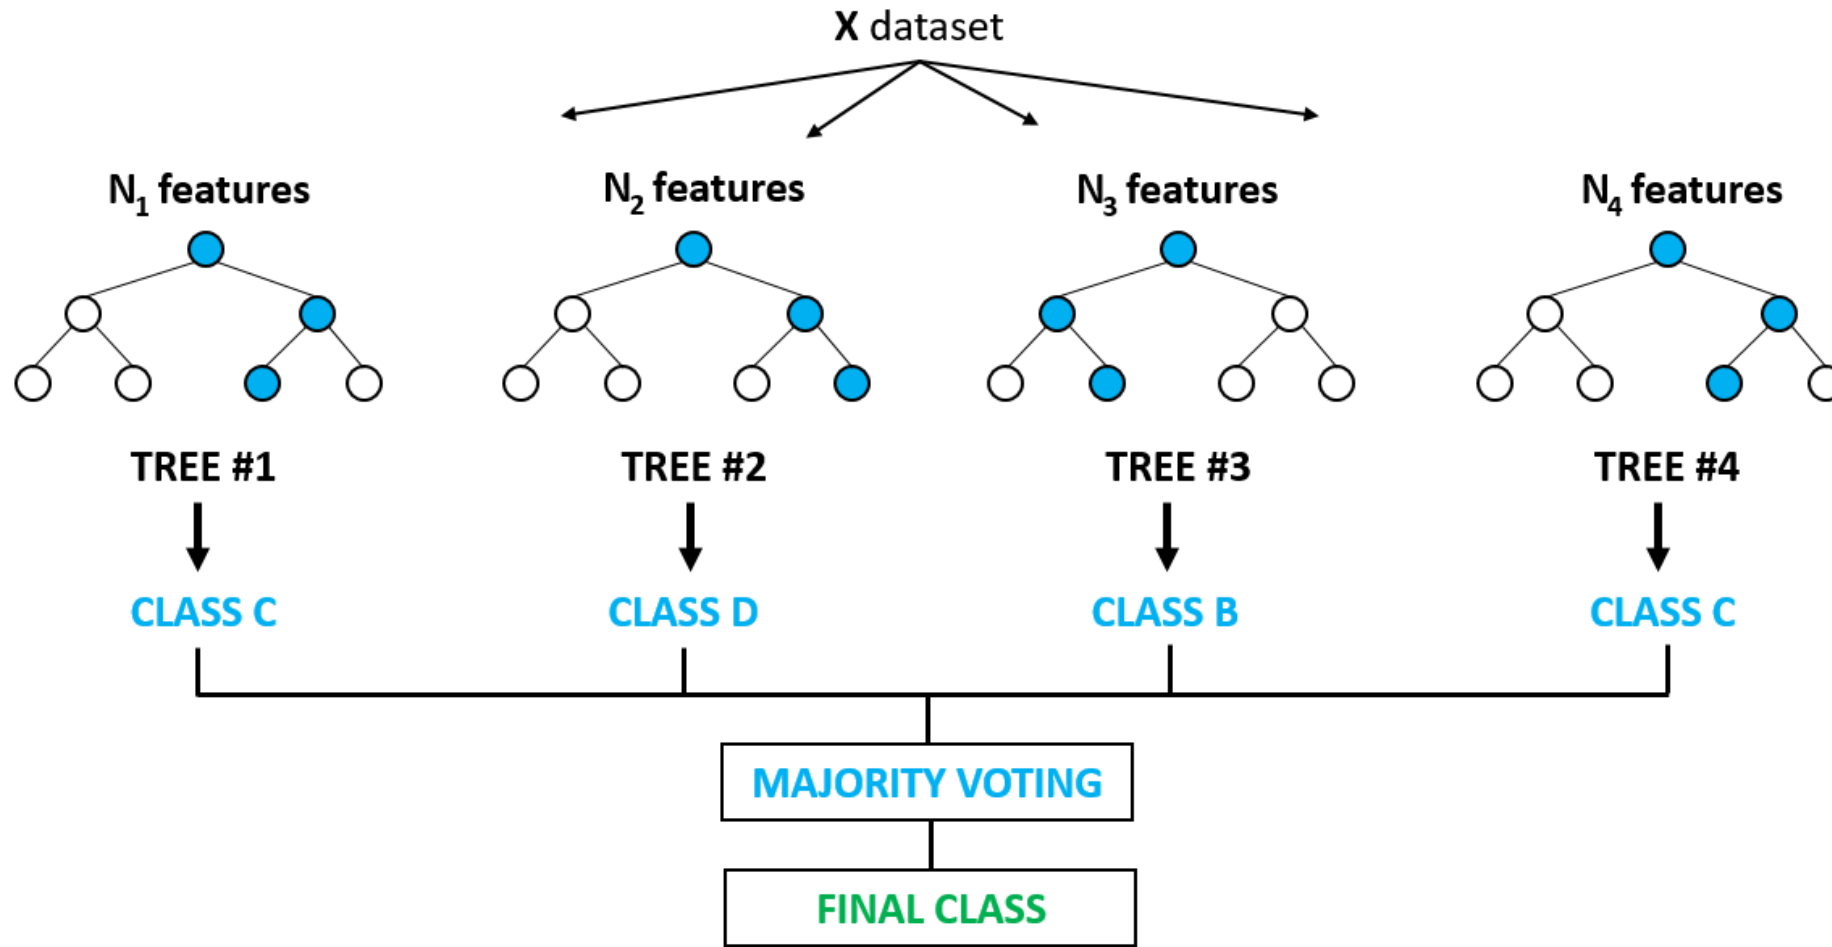
A random forest is constructed by:

1. Choosing the number of decision trees to grow and the number of variables to consider in each tree.
  2. Randomly selecting the rows of the data frame **with replacement**.
  3. Randomly selecting the appropriate number of variables from the data frame.
  4. Building a decision tree on the resulting data set.
  5. Repeating this procedure a large number of times.
  6. A prediction is be made by **majority rule**, i.e. running your new observation through all the trees in the forest and seeing which class is predicted most often.
- The `randomForest::randomForest()` function in R defaults to 500 trees each trained on `sqrt(p)` variables where `p` is the number of predictors in the full data set.

# Ensemble learning: bagging (bootstrap aggregating)



# Ensemble decision making





# randomForest in R

```
library(randomForest)
iris_rf <- randomForest(Species ~ ., iris)
iris_rf

##
## Call:
## randomForest(formula = Species ~ ., data = iris)
##              Type of random forest: classification
##              Number of trees: 500
## No. of variables tried at each split: 2
##
##              OOB estimate of  error rate: 5.33%
## Confusion matrix:
##              setosa versicolor virginica class.error
## setosa           50           0           0          0.00
## versicolor       0           47           3          0.06
## virginica         0           5          45          0.10
```

Using the `randomForest()` function, we can train our ensemble learning using the same formula we passed to `rpart`.

# randomForest in R

```
predict(iris_rf, new_data)
```

```
##      1      2  
## setosa setosa  
## Levels: setosa versicolor virginica
```

Again, we can use the same `new_data` values in the `predict` function as before to reach the same prediction we did with `rpart`.

# Titanic random forest

```
titanic_rf = randomForest(survived ~ sex + age + pclass, titanic_complete)
titanic_rf
```

```
##
## Call:
## randomForest(formula = survived ~ sex + age + pclass, data = titanic_complete)
##           Type of random forest: classification
##           Number of trees: 500
## No. of variables tried at each split: 1
##
##           OOB estimate of  error rate: 20.75%
## Confusion matrix:
##           died survived class.error
## died           567         52  0.08400646
## survived       165        262  0.38641686
```

```
importance(titanic_rf)
```

```
##           MeanDecreaseGini
## sex                122.29587
## age                 21.97414
## pclass              40.82804
```

Note: OOB is short for: [Out-of-bag](#)

For more details on decision trees see Hastie, Tibshirani, and Friedman (2009; section 9.2) and James, Witten, Hastie, et al. (2017; chapter 8).

## References

Hastie, T., R. Tibshirani, and J. Friedman (2009). *The Elements of Statistical Learning*. 2nd ed. Springer Series in Statistics. New York, NY, USA: Springer. URL: <https://web.stanford.edu/~hastie/ElemStatLearn/>.


James, G., D. Witten, T. Hastie, and R. Tibshirani (2017). *An Introduction to Statistical Learning: With Applications in R*. New York: Springer. URL: <https://www-bcf.usc.edu/~gareth/ISL/>.

Jed Wing, M. K. C. from, S. Weston, A. Williams, C. Keefer, A. Engelhardt, T. Cooper, Z. Mayer, B. Kenkel, the R Core Team, M. Benesty, et al. (2018). *caret: Classification and Regression Training*. R package version 6.0-80. URL: <https://CRAN.R-project.org/package=caret>.

Liaw, A. and M. Wiener (2002). "Classification and Regression by randomForest". In: *R News* 2.3, pp. 18-22. URL: <https://CRAN.R-project.org/doc/Rnews/>.

Therneau, T. and B. Atkinson (2018). *rpart: Recursive Partitioning and Regression Trees*. R package version 4.1-13. URL: <https://CRAN.R-project.org/package=rpart>.

## Acknowledgements

Special thanks to [Sarah Romanes](#)  for providing content inspiration from her R Ladies talks on machine learning:

- [sarahromanes.github.io/r-ladies-ML-1](https://sarahromanes.github.io/r-ladies-ML-1)
- [sarahromanes.github.io/r-ladies-ML-2](https://sarahromanes.github.io/r-ladies-ML-2)